

INTRODUCTION TO REGRESSION

-INTRODUCTION TO DATA SCIENCE-

ISL 2.1, 2.2, 3.1, 3.2

Lecturer: Darren Homrighausen, PhD

Preamble:

- Outline the notation for a linear regression model
- Briefly review estimation, prediction, and inference for classical linear regression models
- Give examples of the “classical”, “big data”, and “high dimensional” types of problems

NOTATION RECAP

- We have data $\mathcal{D} = \{(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)\}$
(The **training data**)
- $X \in \mathbb{R}^p$ is a vector of **measurements** for each subject
(Example: $X_i = [1, \text{income}_i, \text{education}_i]^\top$)
- $x \in \mathbb{R}^n$ is a vector of **subjects** for each measurement
(Example: $x_j = [\text{income}_1, \text{income}_2, \dots, \text{income}_n]^\top$)
- X_{ij} is the j^{th} measurement on the i^{th} subject
(Example: $X_{ij} = \text{income}_i$)

NOTATIONAL LANDMINE: representing the j^{th} entry of X

→ A reasonable, but technically sloppy, solution: x_j

A linear model review

A LINEAR MODEL: MULTIPLE REGRESSION

RECALL: For regression, **squared-error** is the usual loss function

→ The Bayes rule w.r.t. this loss function is $f_*(X) = \mathbb{E}Y|X$

Specify the model: $f_*(X) = \beta_0 + X^\top \beta = \beta_0 + \sum_{j=1}^p x_j \beta_j$

(This means that we think the relationship is approximately linear in X)

Then we recover the usual linear regression formulation

$$\mathbb{X} = \begin{bmatrix} x_1 & \cdots & x_p \end{bmatrix} = \begin{bmatrix} X_1^\top \\ X_2^\top \\ \vdots \\ X_n^\top \end{bmatrix} \in \mathbb{R}^{n \times p} \quad \text{and} \quad \mathbb{Y} = \begin{bmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{bmatrix} \in \mathbb{R}^n$$

Commonly, a column $x_0^\top = \underbrace{(1, \dots, 1)}_{n \text{ times}}$ is included

(This encodes an intercept term, with intercept parameter β_0)

We could (should?) seek to find a β such that $\mathbb{Y} \approx \mathbb{X}\beta$

A LINEAR MODEL: POLYNOMIAL EFFECTS

Instead, we may believe

$$f_*(X) = \beta_0 + \sum_{j=1}^p x_j \beta_j + \sum_{j \leq j'}^p x_j x_{j'} \beta_{jj'}$$

Then the **feature** matrix is

$$\mathbb{X} = \begin{bmatrix} x_0 & x_1 & \cdots & x_p & x_1^2 & x_1 x_2 & \cdots & x_p^2 \end{bmatrix}$$

(Here, interpret vector multiplication in the entrywise sense, as in **R**: $x * y$)

This corresponds to the “main and interaction effects” model

Example: Biometrics

EXAMPLE

Suppose we have 4 subjects in an experiment

We record

- BMI
- minutes spent exercising in the last 7 days

We want to predict each subject's resting heart rate

The classic linear model would model the regression function as

$$f_*(X) = \beta_0 + \beta^\top X = \beta_0 + \beta_1 \text{BMI} + \beta_2 \text{exercise}$$

where

$$\begin{aligned} f_*(X) &= \mathbb{E}[\text{resting heart rate} | X] \\ X &= [\text{BMI}, \text{exercise}] \end{aligned}$$

(Note: we could write $f_*(X) = \beta^\top X$ and $X = [1, \text{BMI}, \text{exercise}]$ instead)

EXAMPLE

Under this model, the feature matrix and supervisor vector look like

$$\mathbb{X} = \begin{bmatrix} x_1 & x_2 \end{bmatrix} = \underbrace{\begin{bmatrix} 21 & 92 \\ 17 & 12 \\ 29 & 306 \\ 25 & 53 \end{bmatrix}}_{\text{BMI} \quad \text{exercise}} \in \mathbb{R}^{4 \times 2}$$

and

$$\mathbb{Y} = \begin{bmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_4 \end{bmatrix} = \begin{bmatrix} 72 \\ 47 \\ 82 \\ 64 \end{bmatrix} \in \mathbb{R}^4$$

EXAMPLE

Adding a quadratic polynomial transformation

$$\begin{aligned}f_*(X) &= \beta_0 + \sum_{j=1}^p x_j \beta_j + \sum_{j \leq j'}^p x_j x_{j'} \beta_{jj'} \\&= \beta_0 + \beta_1 \text{BMI} + \beta_2 \text{exercise} + \beta_{11} \text{BMI}^2 + \beta_{22} \text{exercise}^2 \\&\quad + \beta_{12} \text{BMI} \text{exercise}\end{aligned}$$

Under this model, the feature matrix looks like

$$\mathbb{X} = \begin{bmatrix} x_1 & x_2 & x_3 & x_4 & x_5 \end{bmatrix} = \underbrace{\begin{bmatrix} 21 & 92 & 21^2 & 92^2 & 21 * 92 \\ 17 & 12 & 17^2 & 12^2 & 17 * 12 \\ 29 & 306 & 29^2 & 306^2 & 29 * 306 \\ 25 & 53 & 25^2 & 53^2 & 25 * 53 \end{bmatrix}}_{\begin{matrix} \text{BMI} & \text{exercise} & \text{BMI}^2 & \text{exercise}^2 & \text{BMI} * \text{exercise} \end{matrix}}$$

(\mathbb{Y} is the same)

End example

A LINEAR MODEL: ESTIMATING β

In either case, we have a feature matrix \mathbb{X} and supervisor vector \mathbb{Y}

Now, we want to estimate a parameter vector β in the model

$$\mathbb{Y} = \mathbb{X}\beta + \epsilon$$

where $\mathbb{V}\epsilon = \sigma^2$

CLASSICAL LEAST SQUARES: Minimize the training error $\hat{R}(f)$ over all functions $f_\beta(X) = X^\top \beta$

$$\hat{\beta}_{LS} = \underset{\beta}{\operatorname{argmin}} \hat{R}(f_\beta) = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n (Y_i - X_i^\top \beta)^2 = \underset{\beta}{\operatorname{argmin}} \|\mathbb{Y} - \mathbb{X}\beta\|_2^2$$

(Though we write this as equality, there is only a unique solution if $\operatorname{rank}(\mathbb{X}) = p$)

A LINEAR MODEL: PROPERTIES OF $\hat{\beta}_{LS}$

In this case,

$$\hat{f}(X) = X^T \hat{\beta}_{LS} = X^T \underbrace{\mathbb{X}^\dagger}_{\text{rank}(\mathbb{X})=p} Y = X^T (\mathbb{X}^T \mathbb{X})^{-1} \mathbb{X}^T Y$$

(\mathbb{X}^\dagger is a **pseudo inverse**)

The **fitted values** are $\hat{\mathbb{Y}} = \mathbb{X} \hat{\beta}_{LS}$

(Contrary to $\hat{\beta}_{LS}$, the fitted values are always unique)

We can examine the first and second moment properties of $\hat{\beta}_{LS}$

$\mathbb{E} \hat{\beta}_{LS} = \beta$ (unbiased if $f_*(X) = X^T \beta$ is correct model)

$$\mathbb{V} \hat{\beta}_{LS} = \mathbb{X}^\dagger (\mathbb{V} Y) (\mathbb{X}^\dagger)^T = \sigma^2 (\mathbb{X}^T \mathbb{X})^{-1}$$

As $\hat{\beta}_{LS}$ is a fancy average, the **central limit theorem (CLT)** states

$$\hat{\beta}_{LS} \sim N(\beta, \sigma^2 (\mathbb{X}^T \mathbb{X})^{-1})$$

A LINEAR MODEL: INFERENCE USING $\hat{\beta}_{LS}$

Using the CLT result:

$$\hat{\beta}_{LS} \sim N(\beta, \sigma^2(\mathbb{X}^\top \mathbb{X})^{-1})$$

We can test whether $\beta_j = (\text{some value})$ via

$$t_j = \frac{\hat{\beta}_{LS,j} - (\text{some value})}{\sqrt{\mathbb{V}\hat{\beta}_{LS,j}}}$$

where $\mathbb{V}\hat{\beta}_{LS,j}$ is the j^{th} diagonal element of $\sigma^2(\mathbb{X}^\top \mathbb{X})^{-1}$

Under the null hypothesis, $t_j \sim t_{n-p}$

So, large values of $|t_j|$ relative to quantiles of t_{n-p} provides some evidence that $\beta_j \neq (\text{some value})$

End review

TURNING THESE IDEAS INTO PROCEDURES

Each of these methods have parameters to choose:

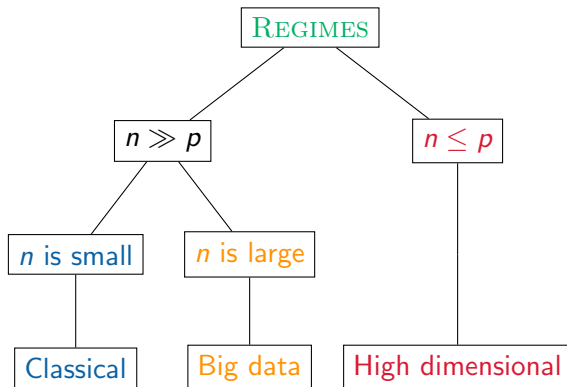
- p could be very large. Do we include all the features?
- If we include some polynomial (or other transformations) terms, should we include all of them?
- Are there other parameters that need to be set in an informed manner?

Additionally, we need to estimate the associated coefficient vector β or whatever

We would like the data to inform these parameters

TURNING THESE IDEAS INTO PROCEDURES

Back to the **three** regimes of interest, assuming $\mathbb{X} \in \mathbb{R}^{n \times p}$



CLASSICAL REGIME

Back to $\hat{\beta}_{LS}$:

The Gauss-Markov theorem assures us that this is the best linear **unbiased** estimator of β

Also, it is the maximum likelihood estimator under the i.i.d. Gaussian model

(Hence, it is asymptotically efficient)

Does that necessarily make it is any good?

CLASSICAL REGIME

Write $\mathbb{X} = UDV^\top$ for the SVD of \mathbb{X}

$$\text{Then } \mathbb{V}\hat{\beta}_{LS} = \sigma^2(\mathbb{X}^\top \mathbb{X})^{-1} = \sigma^2(VD \underbrace{U^\top U}_{=I} DV^\top)^{-1} = \sigma^2 VD^{-2} V^\top$$

(REMINDER: The d_j are the axes lengths of the ellipse of \mathbb{X})

Suppose we are interested in estimating β

Then we want $\mathbb{E}||\hat{\beta}_{LS} - \beta||_2^2$ to be small

(That is, our estimator is close to the true parameter on average)

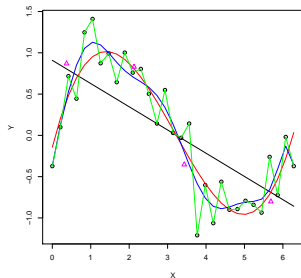
But,

$$\mathbb{E}||\hat{\beta}_{LS} - \beta||_2^2 = \text{trace}(\mathbb{V}\hat{\beta}) = \sigma^2 \sum_{j=1}^p \frac{1}{d_j^2} \quad (1)$$

(Can you show this? Hint: add and subtract $\mathbb{E}\hat{\beta}_{LS}$)

IMPORTANT: Even in the classical regime, we can do arbitrarily badly if $d_p \approx 0$! (An example of this would be “multicollinearity”)

RETURNING TO POLYNOMIAL EXAMPLE: BIAS



Using a Taylor's series, for all x

$$\sin(x) = \sum_{q=0}^{\infty} \frac{(-1)^q x^{2q+1}}{(2q+1)!}$$

Higher order polynomial models will **reduce** the bias part

RETURNING TO POLYNOMIAL EXAMPLE: VARIANCE

The least squares solution is given by solving $\min ||\mathbb{X}\beta - Y||_2^2$

$$\mathbb{X} = \begin{bmatrix} 1 & X_1 & \dots & X_1^{p-1} \\ \vdots & \vdots & & \vdots \\ 1 & X_n & \dots & X_n^{p-1} \end{bmatrix},$$

is the associated feature matrix

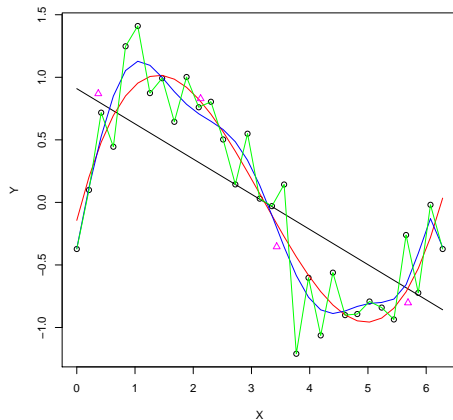
(This is known as the **Vandermonde matrix** in numerical analysis)

This matrix is well known for being numerically unstable due to $d_p \approx 0$

Hence

$$\sum_{j=1}^p \frac{1}{d_j^2} \text{ is huge!}$$

RETURNING TO THE POLYNOMIAL EXAMPLE



CONCLUSION

CONCLUSION: Fitting the full least squares model, even in the classical regime, can lead to poor prediction/estimation performance

In the other regimes, we encounter even more **sinister** problems

BIG DATA REGIME

Big data: Computational/storage complexity scales extremely quickly. This means that procedures that are feasible classically are not for large data sets

EXAMPLE: Fit $\hat{\beta}_{LS}$ with $\mathbb{X} \in \mathbb{R}^{n \times p}$. Next fit $\hat{\beta}_{LS}$ with $\mathbb{X} \in \mathbb{R}^{3n \times 4p}$

The second case will take $\approx (3 * 4^2) = 48$ times longer to compute, as well as ≈ 12 times as much memory!

(In general, the computational complexity scales like np^2)

CONCLUSION

```
p = 300; n = 10000
Y = rnorm(n); X = matrix(rnorm(n*p),nrow=n,ncol=p)
start = proc.time()[3]
out    = lm(Y~.,data=data.frame(X))
end    = proc.time()[3]
smallTime = end - start
```

```
n = nMultiple*n; nMultiple = 3
p = pMultiple*p; pMultiple = 4
Y = rnorm(n); X = matrix(rnorm(n*p),nrow=n,ncol=p)
start = proc.time()[3]
out    = lm(Y~.,data=data.frame(X))
end    = proc.time()[3]
bigTime = end - start
> print(bigTime/smallTime)
  elapsed
38.61458
> print(nMultiple*pMultiple**2)
[1] 48
```

TREATMENT IN PRACTICE

Depending on the data and the desired method, we could:

- Combine randomized projections together with in-memory procedures
(EXAMPLE: We can randomly subsample observations and then load into memory)
- Use (stochastic) gradient descent
(We will return to this later)
- Leverage an iterative implementation for exact computation
(EXAMPLE: `biglm` in R)
- Break the computations down into small bits and distribute these to different cores/processors/nodes
(This is like the map-reduce paradigm)

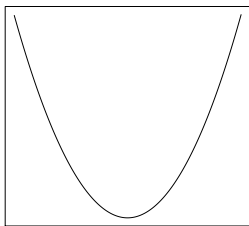
HIGH DIMENSIONAL REGIME

High dimensional: These problems tend to have many of the computational problems as **Big data**, as well as a **rank problem**:

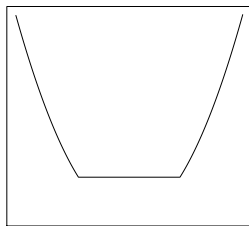
Suppose $\mathbb{X} \in \mathbb{R}^{n \times p}$ and $p > n$

Then $\text{rank}(\mathbb{X}) = n$ and the equation $\mathbb{X}\hat{\beta} = Y$:

- can be solved *exactly* (that is; the training error is 0)
- has an infinite number of solutions



$n > p$



$n < p$

Postamble:

- Outline the notation for a linear regression model
(Write $\mathbb{Y} = \mathbb{X}\beta + \epsilon$)
- Briefly review estimation, prediction, and inference for linear regression models
(Least squares is same as minimizing training error with squared error loss)
- Define and give examples of the “classical”, “big data”, and “high dimensional” types of problems

