

# SUPPORT VECTOR MACHINES 2

-INTRODUCTION TO DATA SCIENCE-

Lecturer: Darren Homrighausen, PhD

# Preamble:

- We will generalize the Support Vector Classifier into the Support Vector Machine by leveraging “kernelization”
- Kernelization is the idea that we can replace inner products of observations ( $\langle X_i, X_{i'} \rangle = X_i^\top X_{i'}$ ) with kernel evaluations ( $k(X_i, X_{i'})$ )
- This allows us to do dimension **expansion** without increasing the computational burden
- There is still a danger of overfitting with kernel methods, so we must regularize. Hence, we show that SVMs can be written as a penalized loss method, just like the logistic elastic net

# OPTIMAL SEPARATING HYPERPLANE

**REMINDER:** The **optimal separating hyperplane** is produced by maximizing

$$\sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{k=1}^n \alpha_i \alpha_k Y_i Y_k \mathbf{x}_i^\top \mathbf{x}_k$$

subject to  $\alpha_i \geq 0$

A similar result holds after the introduction of slack variables (e.g. **support vector classifiers**). In fact, the only difference is  $\alpha_i \leq \lambda$  for each  $i$

**IMPORTANT:** The features only enter via

$$\mathbf{x}^\top \mathbf{x}' = \langle \mathbf{x}, \mathbf{x}' \rangle$$

# (KERNEL) RIDGE REGRESSION

REMINDER: Suppose we want to predict at  $X$ , then

$$\hat{f}(X) = X^\top \hat{\beta}_{\text{ridge}}(\lambda) = X^\top \mathbb{X}^\top (\mathbb{X} \mathbb{X}^\top + \lambda I)^{-1} Y$$

Also,

$$\mathbb{X} \mathbb{X}^\top = \begin{bmatrix} \langle X_1, X_1 \rangle & \langle X_1, X_2 \rangle & \cdots & \langle X_1, X_n \rangle \\ & \vdots & & \\ \langle X_n, X_1 \rangle & \langle X_n, X_2 \rangle & \cdots & \langle X_n, X_n \rangle \end{bmatrix}$$

and

$$X^\top \mathbb{X}^\top = [\langle X, X_1 \rangle, \langle X, X_2 \rangle, \dots, \langle X, X_n \rangle]$$

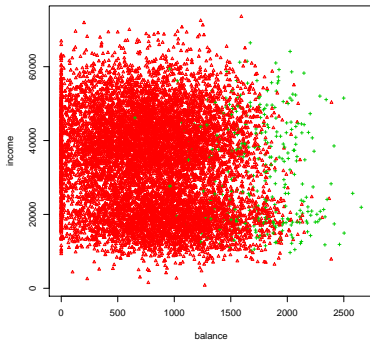
Again, we have the features enter only as

$$\langle X, X' \rangle = X^\top X'$$

# LOGISTIC REGRESSION: TRANSFORMATIONS

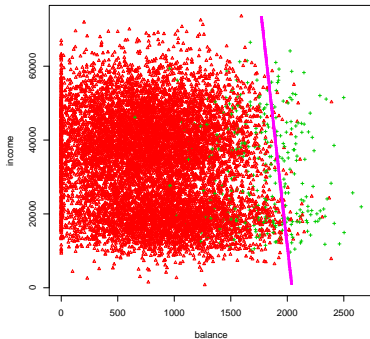
Let's look at the **default** data in ISL

In particular, we will look at **default** status as a function of **balance** and **income**



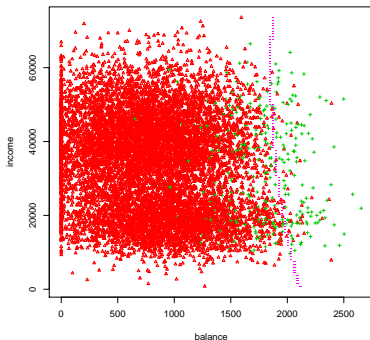
# LOGISTIC REGRESSION: TRANSFORMATIONS

```
out.glm = glm(default~balance + income,family='binomial')
```



# LOGISTIC REGRESSION: TRANSFORMATIONS

```
out.glm = glm(default~balance + income +  
              I(income^2),family='binomial')
```



**CONCLUSION:** Linear rules in a transformed space can have nonlinear decisions in original features

# LOGISTIC REGRESSION: TRANSFORMATIONS

REMINDER: The logistic model: untransformed

$$\begin{aligned}\text{logit}(\mathbb{P}(Y = 1|X)) &= \beta_0 + \beta^\top X \\ &= \beta_0 + \beta_1 \text{balance} + \beta_2 \text{income}\end{aligned}$$

The decision boundary is the hyperplane  $\{X : \beta_0 + \beta^\top X = 0\}$

This is **linear** in the feature space



# LOGISTIC REGRESSION: TRANSFORMATIONS

Adding the polynomial transformation:

$$\Phi(X) = (\phi_1(X), \phi_2(X), \phi_3(X)) = (x_1, x_2, x_2^2)$$

$$\begin{aligned}\text{logit}(\mathbb{P}(Y = 1|X)) &= \beta_0 + \beta^\top \Phi(X) \\ &= \beta_0 + \beta_1 \text{balance} + \beta_2 \text{income} + \beta_3 \text{income}^2\end{aligned}$$

Decision boundary is still a hyperplane  $\{X : \beta_0 + \beta^\top \Phi(X) = 0\}$

This is **nonlinear** in the original  $x_1, x_2$  feature space, but is linear in terms of  $\Phi(X)$ !

# LOGISTIC REGRESSION: TRANSFORMATIONS

Of course, as we include more transformations,

- We need to choose the transformations **manually**
- **Computations** can become difficult if we aren't careful
- We need to **regularize** to prevent overfitting

Can we form them in an automated fashion?

# Kernel Methods

# KERNEL: EXAMPLE

Back to polynomial terms/interactions  $\Phi(X) = (x_1, x_2, x_2^2)$ :

What if instead we could form a (kernel) function that produces these polynomial terms **automatically**?

WE CAN!

$$\rightarrow \text{Form } k(X, X') = (X^\top X' + 1)^d$$

This kernel **implicitly** forms all polynomials/interactions up to degree  $d$

# KERNEL: EXAMPLE

**EXAMPLE:** Let  $d = p = 2$  and  $u, v \in R^2$  be two vectors

$$\begin{aligned}k(u, v) &= 1 + 2u_1v_1 + 2u_2v_2 + u_1^2v_1^2 + u_2^2v_2^2 + 2u_1u_2v_1v_2 \\&= \sum_{k=1}^M \Phi_k(u)\Phi_k(v) \quad (M = 6) \\&= \Phi(u)^\top \Phi(v) \\&= \langle \Phi(u), \Phi(v) \rangle\end{aligned}$$

where

$$\Phi(v)^\top = [1, \sqrt{2}v_1, \sqrt{2}v_2, v_1^2, v_2^2, \sqrt{2}v_1v_2]$$

**IMPORTANT:** These equalities are **everything** that makes kernelization work!

# KERNEL: CONCLUSION

Let's recap:

$$\begin{aligned}k(u, v) &= 1 + 2u_1v_1 + 2u_2v_2 + u_1^2v_1^2 + u_2^2v_2^2 + 2u_1u_2v_1v_2 \\ &= \langle \Phi(u), \Phi(v) \rangle = \Phi(u)^\top \Phi(v)\end{aligned}$$

where

$$\Phi(v)^\top = [1, \sqrt{2}v_1, \sqrt{2}v_2, v_1^2, v_2^2, \sqrt{2}v_1v_2]$$

- Some methods only involve features via inner products  
 $X^\top X' = \langle X, X' \rangle$   
(We've explicitly seen two: ridge regression and support vector classifiers)
- If we make transformations of  $X$  to  $\Phi(X)$ , the procedure depends on  $\Phi(X)^\top \Phi(X') = \langle \Phi(X), \Phi(X') \rangle$
- **CRUCIAL:** We can compute this inner product via the kernel:

$$k(X, X') = \langle \Phi(X), \Phi(X') \rangle$$

## KERNEL: CONCLUSION

Instead of creating a very high dimensional object via transformations, choose a kernel  $k$

Now, the only thing left to do is form the **outer product** of kernel evaluations

$$\mathbb{K} = [k(X_i, X_{i'})]_{1 \leq i, i' \leq n}$$

```
X = c(1,2,3)
k = function(x,y){ return(x + y + x*y)}
> outer(X,X,k)
      [,1] [,2] [,3]
[1,]    3    5    7
[2,]    5    8   11
[3,]    7   11   15
```

# (Kernel) SVMs



# KERNEL SVM

RECALL:

$$\frac{1}{2} \|\beta\|_2^2 - \sum_{i=1}^n \alpha_i [Y_i (X_i^\top \beta + \beta_0) - 1]$$

Derivatives with respect to  $\beta$  and  $\beta_0$  imply:

- $\beta = \sum_{i=1}^n \alpha_i Y_i X_i$
- $0 = \sum_{i=1}^n \alpha_i Y_i$

Write the solution function

$$f(X) = \beta_0 + \beta^\top X = \beta_0 + \sum_{i=1}^n \alpha_i Y_i X_i^\top X$$

Kernelize the SVC  $\Rightarrow$  support vector machine (SVM):

$$f(X) = \beta_0 + \sum_{i=1}^n \alpha_i Y_i k(X_i, X)$$



# GENERAL KERNEL MACHINES



We can write the eigenvalue expansion of  $k$  as

$$k(u, v) = \sum_{j=1}^{\infty} \theta_j \phi_j(u) \phi_j(v)$$

(This is called **Mercer's theorem**, and such a  $k$  is called a **Mercer kernel**)

Replacing inner products with kernel evaluations is equivalent to performing the unkernelized method in the space given by the **eigenfunctions** of  $k$  with **feature map**  $\Phi = [\phi_1, \dots, \phi_p, \dots]$

## POLYNOMIAL EXAMPLE:

$$\begin{aligned} k(u, v) &= 1 + 2u_1v_1 + 2u_2v_2 + u_1^2v_1^2 + u_2^2v_2^2 + 2u_1u_2v_1v_2 \\ &= \sum_{k=1}^M \Phi_k(u) \Phi_k(v), \end{aligned}$$

where:  $\Phi(v)^\top = [1, \sqrt{2}v_1, \sqrt{2}v_2, v_1^2, v_2^2, \sqrt{2}v_1v_2]$

# KERNEL SVMs

Hence (and luckily) specifying  $\Phi$  itself unnecessary,

We need only define the **kernel** that is symmetric, positive definite

Some common choices for SVMs:

- **POLYNOMIAL:**  $k(X, X') = (1 + X^\top X')^d$
- **RADIAL BASIS:**  $k(X, X') = e^{-\tau \|X - X'\|_b^b}$

(For example,  $b = 2$  and  $\tau = \sigma^{-2}/2$  is (proportional to) the Gaussian density)

# KERNEL SVMs: SUMMARY

Reminder: the solution form for SVM is

$$\beta = \sum_{i=1}^n \alpha_i Y_i X_i$$

Kernelized, this is

$$\beta = \sum_{i=1}^n \alpha_i Y_i \Phi(X_i)$$

Therefore, the induced hyperplane is:

$$\begin{aligned} f(X) &= \Phi(X)^\top \beta + \beta_0 = \sum_{i=1}^n \alpha_i Y_i \langle \Phi(X), \Phi(X_i) \rangle + \beta_0 \\ &= \sum_{i=1}^n \alpha_i Y_i k(X, X_i) + \beta_0 \end{aligned}$$

The final classification is still  $\hat{g}(X) = \text{sgn}(\hat{f}(X))$

# SVMs via penalization

# SVMs VIA PENALIZATION

**NOTE:** SVMs can be derived from **penalized loss** methods

The support vector classifier optimization problem:

$$\min_{\beta_0, \beta, \xi} \frac{1}{2} \|\beta\|_2^2 + \lambda \sum \xi_i \quad \text{subject to}$$
$$Y_i f(X_i) \geq 1 - \xi_i, \xi_i \geq 0, \text{ for each } i$$

Consider the alternative optimization problem:

$$\min_{\beta, \beta_0} \sum_{i=1}^n [1 - Y_i f(X_i)]_+ + \tau \|\beta\|_2^2$$

These optimization problems are the same!

(With the relation:  $2\lambda = 1/\tau$ )

# SVMs VIA PENALIZATION

The **loss** part is the **hinge loss function**

$$\ell(f(X), Y) = [1 - Yf(X)]_+$$

The hinge loss approximates the zero-one loss function underlying classification

It has one major advantage, however: **convexity**

# SURROGATE LOSSES

Looking at

$$\min_{\beta, \beta_0} \sum_{i=1}^n [1 - Y_i f(X_i)]_+ + \tau \|\beta\|_2^2$$

It is tempting to minimize (analogous to linear regression)

$$\sum_{i=1}^n \mathbf{1}(Y_i \neq g(X_i)) + \tau \|\beta\|_2^2$$

However, this is **nonconvex**



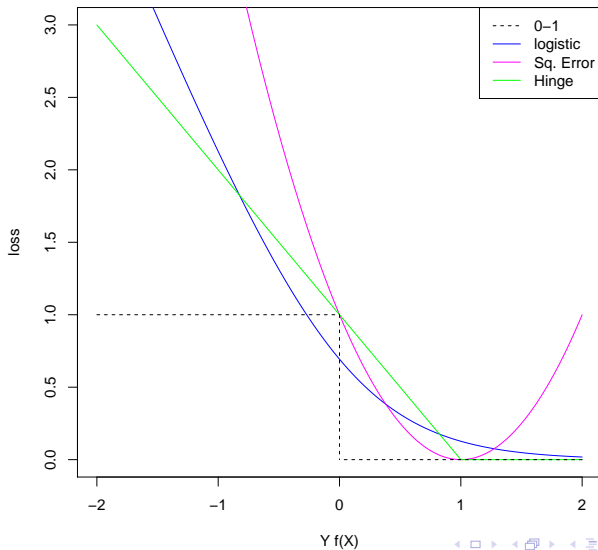
# SURROGATE LOSSES

**IDEA:** We can use a **surrogate** loss that mimics this function while still being convex

It turns out we have already done that! (two times)

- **HINGE:**  $[1 - Yf(X)]_+$
- **LOGISTIC:**  $\log(1 + e^{-Yf(X)})$

# COMPARING LOSS FUNCTIONS



# SVMs IN PRACTICE

**GENERAL FUNCTIONS:** The basic SVM functions are in the C++ library **libsvm**

**R PACKAGE:** The **R** package **e1071** calls **libsvm**

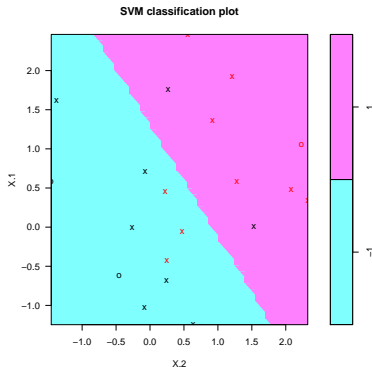
**PATH ALGORITHM:** **svmpath**

For a nice comparison of these approaches, see “Support vector machines in **R**”

(<http://www.jstatsoft.org/v15/i09/paper>)

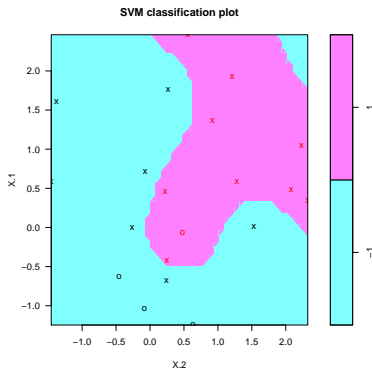
# SVM EXAMPLE

```
tune.out = tune(svm,Y~.,data=dat,kernel="linear",  
  ranges=list(cost=c(0.001, 0.01, 0.1, 1,5,10,100)))
```



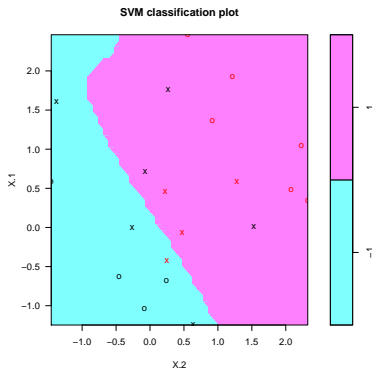
# SVM EXAMPLE

```
tune.out = tune(svm,Y~.,data=dat,kernel="radial",  
  gamma=c(1,2),  
  ranges=list(cost=c(0.001, 0.01, 0.1, 1,5,10,100)))
```

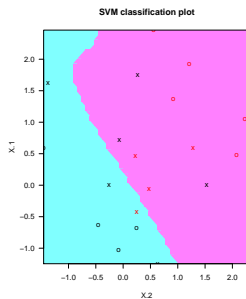
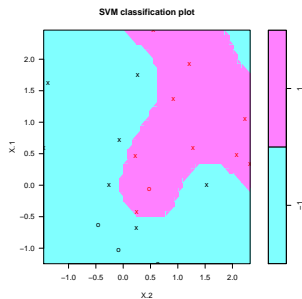
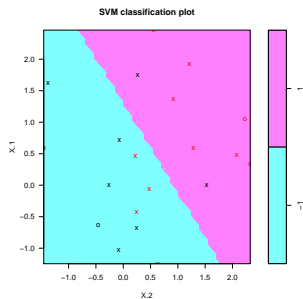


# SVM EXAMPLE

```
tune.out = tune(svm,Y~.,data=dat,kernel="polynomial",  
  degree=c(3,5,10),  
  ranges=list(cost=c(0.001, 0.01, 0.1, 1,5,10,100)))
```



# SVM EXAMPLE



# Multiclass SVMs



# MULTICLASS SVMs

Sometimes, it becomes necessary to do multiclass classification

There are two main approaches:

- One-versus-one
- One-versus-all

# MULTICLASS SVMs: ONE-VERSUS-ONE

Here, for  $G$  possible classes, we run  $G(G - 1)/2$  possible pairwise classifications

For a given test point  $X$ , we find  $\hat{g}_k(X)$  for  $k = 1, \dots, G(G - 1)/2$  fits

The result is a vector  $\hat{G} \in \mathbb{R}^G$  with the total number of times  $X$  was assigned to each class

We report  $\hat{g}(X) = \arg \max_g \hat{G}$

This approach leverages all the class information, but can be **REALLY** slow

(It does have the advantage of only using at any one time the training observations  $i$  such that  $Y_i$  has either of two classes we are considering)

# MULTICLASS SVMs: ONE-VERSUS-ALL

Here, we fit only  $G$  SVMs by respectively collapsing over all size  $G - 1$  subsets of  $\{1, \dots, G\}$

Take all  $\hat{f}_g(X)$  for  $g = 1, \dots, G$ , where class  $g$  is coded 1 and “the rest” is coded -1

Assign  $\hat{g}(X) = \arg \max_g \hat{f}_g(X)$

# Postamble:

- We will generalize the Support Vector Classifier into the Support Vector Machine by leveraging “kernelization”

(Reminder: we saw this before in ridge regression)

- Kernelization is the idea that we can replace inner products of observations ( $\langle X_i, X_{i'} \rangle = X_i^\top X_{i'}$ ) with kernel evaluations ( $k(X_i, X_{i'})$ )

(This can be seen in the form of the solution

$$\hat{f}(X) = \sum_{i=1}^n \alpha_i Y_i X^\top X_i \rightarrow \sum_{i=1}^n \alpha_i Y_i k(X, X_i))$$

- This allows us to do dimension **expansion** without increasing the computational burden
- There is still a danger of overfitting with kernel methods, so we must regularize. Hence, we show that SVMs can be written as a penalized loss method, just like the logistic elastic net

(This is the “hinge loss function”)