

# BOOSTING 3: IMPLEMENTATIONS

-INTRODUCTION TO DATA SCIENCE-

Lecturer: Darren Homrighausen, PhD

# OUTLINE

Now we will discuss two current, popular algorithms and their **R** implementations

- **GBM**
- **XGBoost**

# GBM

# GRADIENT BOOSTING MACHINES (GBM)

**RECALL:** Boosting uses **forward stepwise** minimization of a loss function to build an additive model

**GBM** takes this idea and implements this for

- lots of loss functions
- adds subsampling
- includes methods for choosing  $B$
- reports variable importance measures

# GBM: LOSS FUNCTIONS

- **gaussian**: squared error
- **laplace**: absolute value  
(The Bayes' rule for this loss is the  $\text{median}(Y|X)$ . Hence, it is more robust than for squared error loss)
- **bernoulli**: the Bernoulli family with logistic link
- **adaboost**: exponential
- **multinomial**: more than two classes
- **poisson**: Count data  
(Useful if the supervisor is smallish number of counts (e.g. less than 20))
- **coxph**: For right censored, survival data

# GBM: SUBSAMPLING

It has been noted that boosting performance can be improved if, at each step  $b$ , a new random subsample of observations is used

This is known as **stochastic** gradient boosting

This has two possible benefits

- Reduces computations/storage  
(But increases read/write time)
- Can **improve** performance

(This is the **bag.fraction** parameter)

# GBM: SUBSAMPLING

The improvement due to subsampling suggests the usual ‘variance reduction through lowering covariance’ interpretation

The effect is complicated, though as subsampling

- increases the **variance** of each term in the sum  
(Due to fewer terms training each base learner)
- decreases the **covariance** between each term in the sum  
(Forcing the base learner to focus on different observations)

Another way of viewing this is that subsampling **regularizes** the boosting procedure

# GBM: CHOOSING $B$

There are three built in methods:

- **INDEPENDENT TEST SET:** using the `nTrain` parameter to say 'use only this amount of data for training'  
(Be sure to uniformly permute your data set first.)
- **OUT-OF-BAG (OOB) ESTIMATION:** If `bag.fraction` is  $> 0$ , then `gbm` use OOB at each iteration to find a good  $B$

(Note: OOB tends to select a too-small  $B$ )

- **$K$ -FOLD CROSS VALIDATION (CV):** It will fit `cv.folds+1` models

(The '+1' is the fit on all the data that is reported)

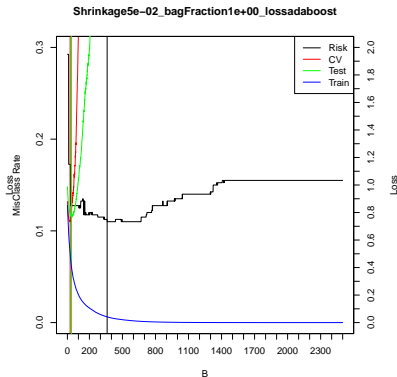
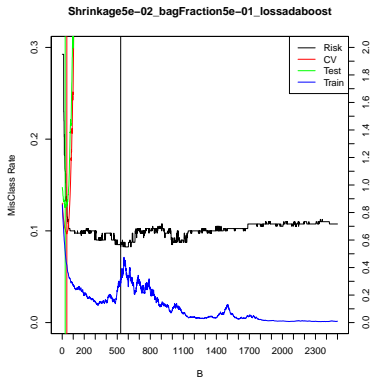


# Simulation

# GBM: BAG FRACTION

Let's look at a simulation for

- adaboost
- bag fraction = 0.5 vs. 1.0
- Test, CV, and training estimates of the risk as well as the actual risk





Bag fraction of 0.5 improves the best risk

# GBM: LOSS FUNCTION


Let's look at a simulation for

- adaboost vs. Bernoulli loss
- bag fraction = 0.5
- Test, CV, and training estimates of the risk as well as the actual risk



Bernoulli loss: CV-minimum  $B \approx$  Risk-minimum  $B$

# COMPARING LOSS FUNCTIONS



**FIGURE:** Here, I've rescaled the logistic loss so that it is easier to compare to Adaboost

**REMINDER:** The boosting classifier is of the form  $\text{sgn}(f(X))$

# GBM: VARIABLE IMPORTANCE MEASURE

For tree-based methods, there are **variable importance measures**:

**RELATIVE.INFLUENCE**: For each feature  $X_j$  and tree  $T_b$

$$\text{Influence}_j(T_b) = \sum_{\text{Split on } X_j} (\text{Reduction in loss})^2$$

This is aggregated to form

$$\text{Influence}_j = \frac{1}{B} \sum_{b=1}^B \text{Influence}_j(T_b)$$

(There is also **permutation.test.gbm**, but it is currently labeled experimental)

# GBM: VARIABLE IMPORTANCE MEASURE

Example using the spam data:

```
gbm.sum = summary(boost.out)
head(gbm.sum)
```

	var	rel.inf
punc_exclam	punc_exclam	22.564132
punc_dollar	punc_dollar	20.547832
remove	remove	11.900176
hp	hp	7.977975
free	free	6.563215
capAvg	capAvg	5.019234

# PARTIAL DEPENDENCE PLOT

GBM provides additional plots for the effect of each feature on the final prediction

The plots are analogous to interpreting a coefficient in multiple regression:

$$f(X) = \sum_{j=1}^p \beta_j x_j$$

The coefficient  $\beta_j$  is the difference in the mean of  $Y$  for a 1 unit change in  $x_j$  **given all the other features are in the model and are held constant**


Boosting fits a nonparametric model, so the interpretation is more complicated

The idea is that we can 'integrate' out the effect of the other features in order to 'hold them constant'

# PARTIAL DEPENDENCE PLOT

Let's look at the the plot for the most important feature

```
most.influential = which(names(X)%in%gbm.sum[1:1,1])  
plot(boost.out,i.var=most.influential)
```






# PARTIAL DEPENDENCE PLOT

Let's look at the the plot for the most important feature

```
most.influential = which(names(X)%in%gbm.sum[1:2,1])  
plot(boost.out,i.var=most.influential)
```



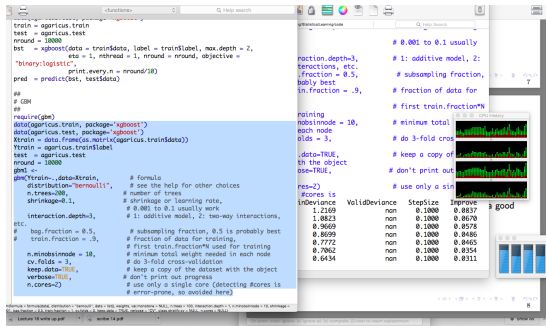
# GBM: SAMPLE CODE

```
gbm(Ytrain~.,data=Xtrain,  
    distribution="bernoulli",  
    n.trees=500,  
    shrinkage=0.01,  
    interaction.depth=3,  
    bag.fraction = 0.5,  
    n.minobsinnode = 10,  
    cv.folds = 3,  
    keep.data=TRUE,  
    verbose=TRUE,  
    n.cores=2)
```


Keep adding trees with **gbm.more**

( If this is taking too long, increase the learning rate, **shrinkage**)

# GBM: FIGURES



# DISTRIBUTED COMPUTING HIERARCHY



**EXAMPLE:** A server might have


- 64 nodes
- 2 processors per node
- 16 cores per processor
- **hyper threading**

The goal is to somehow allocate a **job** so that these resources are used efficiently

Jobs are composed of **threads**, which are specific computations

# HYPERTHREADING

Developed by Intel, Hyperthreading allows for each core to pretend to be two cores



This works by trading off computation and read-time for each core

# BOOSTING: LEARNING SLOW

It is best to set the **learning rate** at a small number.

This is usually calibrated by the computational demands of the problem.

A good strategy is to pick a number, say .001

Run with **n.trees** relatively small and see how long it takes

Keep adding trees with **gbm.more**. If this is taking too long, increase the learning rate

# XGBoost

# XGBOOST

This stands for:

EXTREME GRADIENT BOOSTING

It has some advances related to **gbm**



# XGBOOST: ADVANCES

- **SPARSE MATRICES:** Can use sparse matrices as inputs  
(In fact, it has its own matrix-like data structure that is recommended)
- **OPENMP:** Incorporates OpenMP on Windows/Linux  
(OpenMP is a **message passing** parallelization paradigm for shared memory parallel programming)
- **LOSS FUNCTIONS:** You can specify your own loss/evaluation functions  
(You need to use **xgb.train** for this)